

Design (and Implementation) of PDMP Monte Carlo Methods

BAMC 2022 Minisymposium

Nonreversible Processes: Analysis and Computations

Sam Power, 13 Apr 2022

Problem Setting

Problem Setting

- goal: generate samples from $\pi(x) = \exp(-U(x))/Z$

Problem Setting

- goal: generate samples from $\pi(x) = \exp(-U(x))/Z$
 - easy to evaluate U , Z not so much

Problem Setting

- goal: generate samples from $\pi(x) = \exp(-U(x))/Z$
 - easy to evaluate U , Z not so much
 - if $\dim x \gg 1$, direct strategies unlikely to work

Problem Setting

- goal: generate samples from $\pi(x) = \exp(-U(x))/Z$
 - easy to evaluate U , Z not so much
 - if $\dim x \gg 1$, direct strategies unlikely to work
- strategy: iterative method

Problem Setting

- goal: generate samples from $\pi(x) = \exp(-U(x))/Z$
 - easy to evaluate U , Z not so much
 - if $\dim x \gg 1$, direct strategies unlikely to work
- strategy: iterative method
 - generate a sequence $(x_t)_{t \geq 0}$ such that $\text{Law}(x_t) \xrightarrow{t \rightarrow \infty} \pi$

Problem Setting

- goal: generate samples from $\pi(x) = \exp(-U(x))/Z$
 - easy to evaluate U , Z not so much
 - if $\dim x \gg 1$, direct strategies unlikely to work
- strategy: iterative method
 - generate a sequence $(x_t)_{t \geq 0}$ such that $\text{Law}(x_t) \xrightarrow{t \rightarrow \infty} \pi$
 - for simplicity: x_t memoryless, stationary (“MCMC”)

Dynamical Systems for Sampling

(focus on *local* dynamics, in *continuous time* (makes life easier))

Dynamical Systems for Sampling

(focus on *local* dynamics, in *continuous time* (makes life easier))

- what can dynamics involve?

Dynamical Systems for Sampling

(focus on *local* dynamics, in *continuous time* (makes life easier))

- what can dynamics involve?
 1. following a vector field

Dynamical Systems for Sampling

(focus on *local* dynamics, in *continuous time* (makes life easier))

- what can dynamics involve?
 1. following a vector field
 2. diffusing locally according to some metric

Dynamical Systems for Sampling

(focus on *local* dynamics, in *continuous time* (makes life easier))

- what can dynamics involve?
 1. following a vector field
 2. diffusing locally according to some metric
 3. ‘structured’ jumps

Dynamical Systems for Sampling

(focus on *local* dynamics, in *continuous time* (makes life easier))

- what can dynamics involve?
 1. following a vector field
 2. diffusing locally according to some metric
 3. ‘structured’ jumps
- have some fun: hybrid strategies

Dynamical Systems for Sampling

(focus on *local* dynamics, in *continuous time* (makes life easier))

- what can dynamics involve?
 1. following a vector field
 2. diffusing locally according to some metric
 3. ‘structured’ jumps
- have some fun: hybrid strategies
 - $1 = ?$, $1 + 2 = ?$, $1 + 3 = ?$, $1 + 2 + 3 = ?$


Dynamical Systems for Sampling

(focus on *local* dynamics, in *continuous time* (makes life easier))

- what can dynamics involve?
 1. following a vector field
 2. diffusing locally according to some metric
 3. ‘structured’ jumps
- have some fun: hybrid strategies
 - $1 = ?$, $1 + 2 = ?$, $1 + 3 = ?$, $1 + 2 + 3 = ?$
 - today: PDMPs = $1 + 3$

Primer on PDMPs

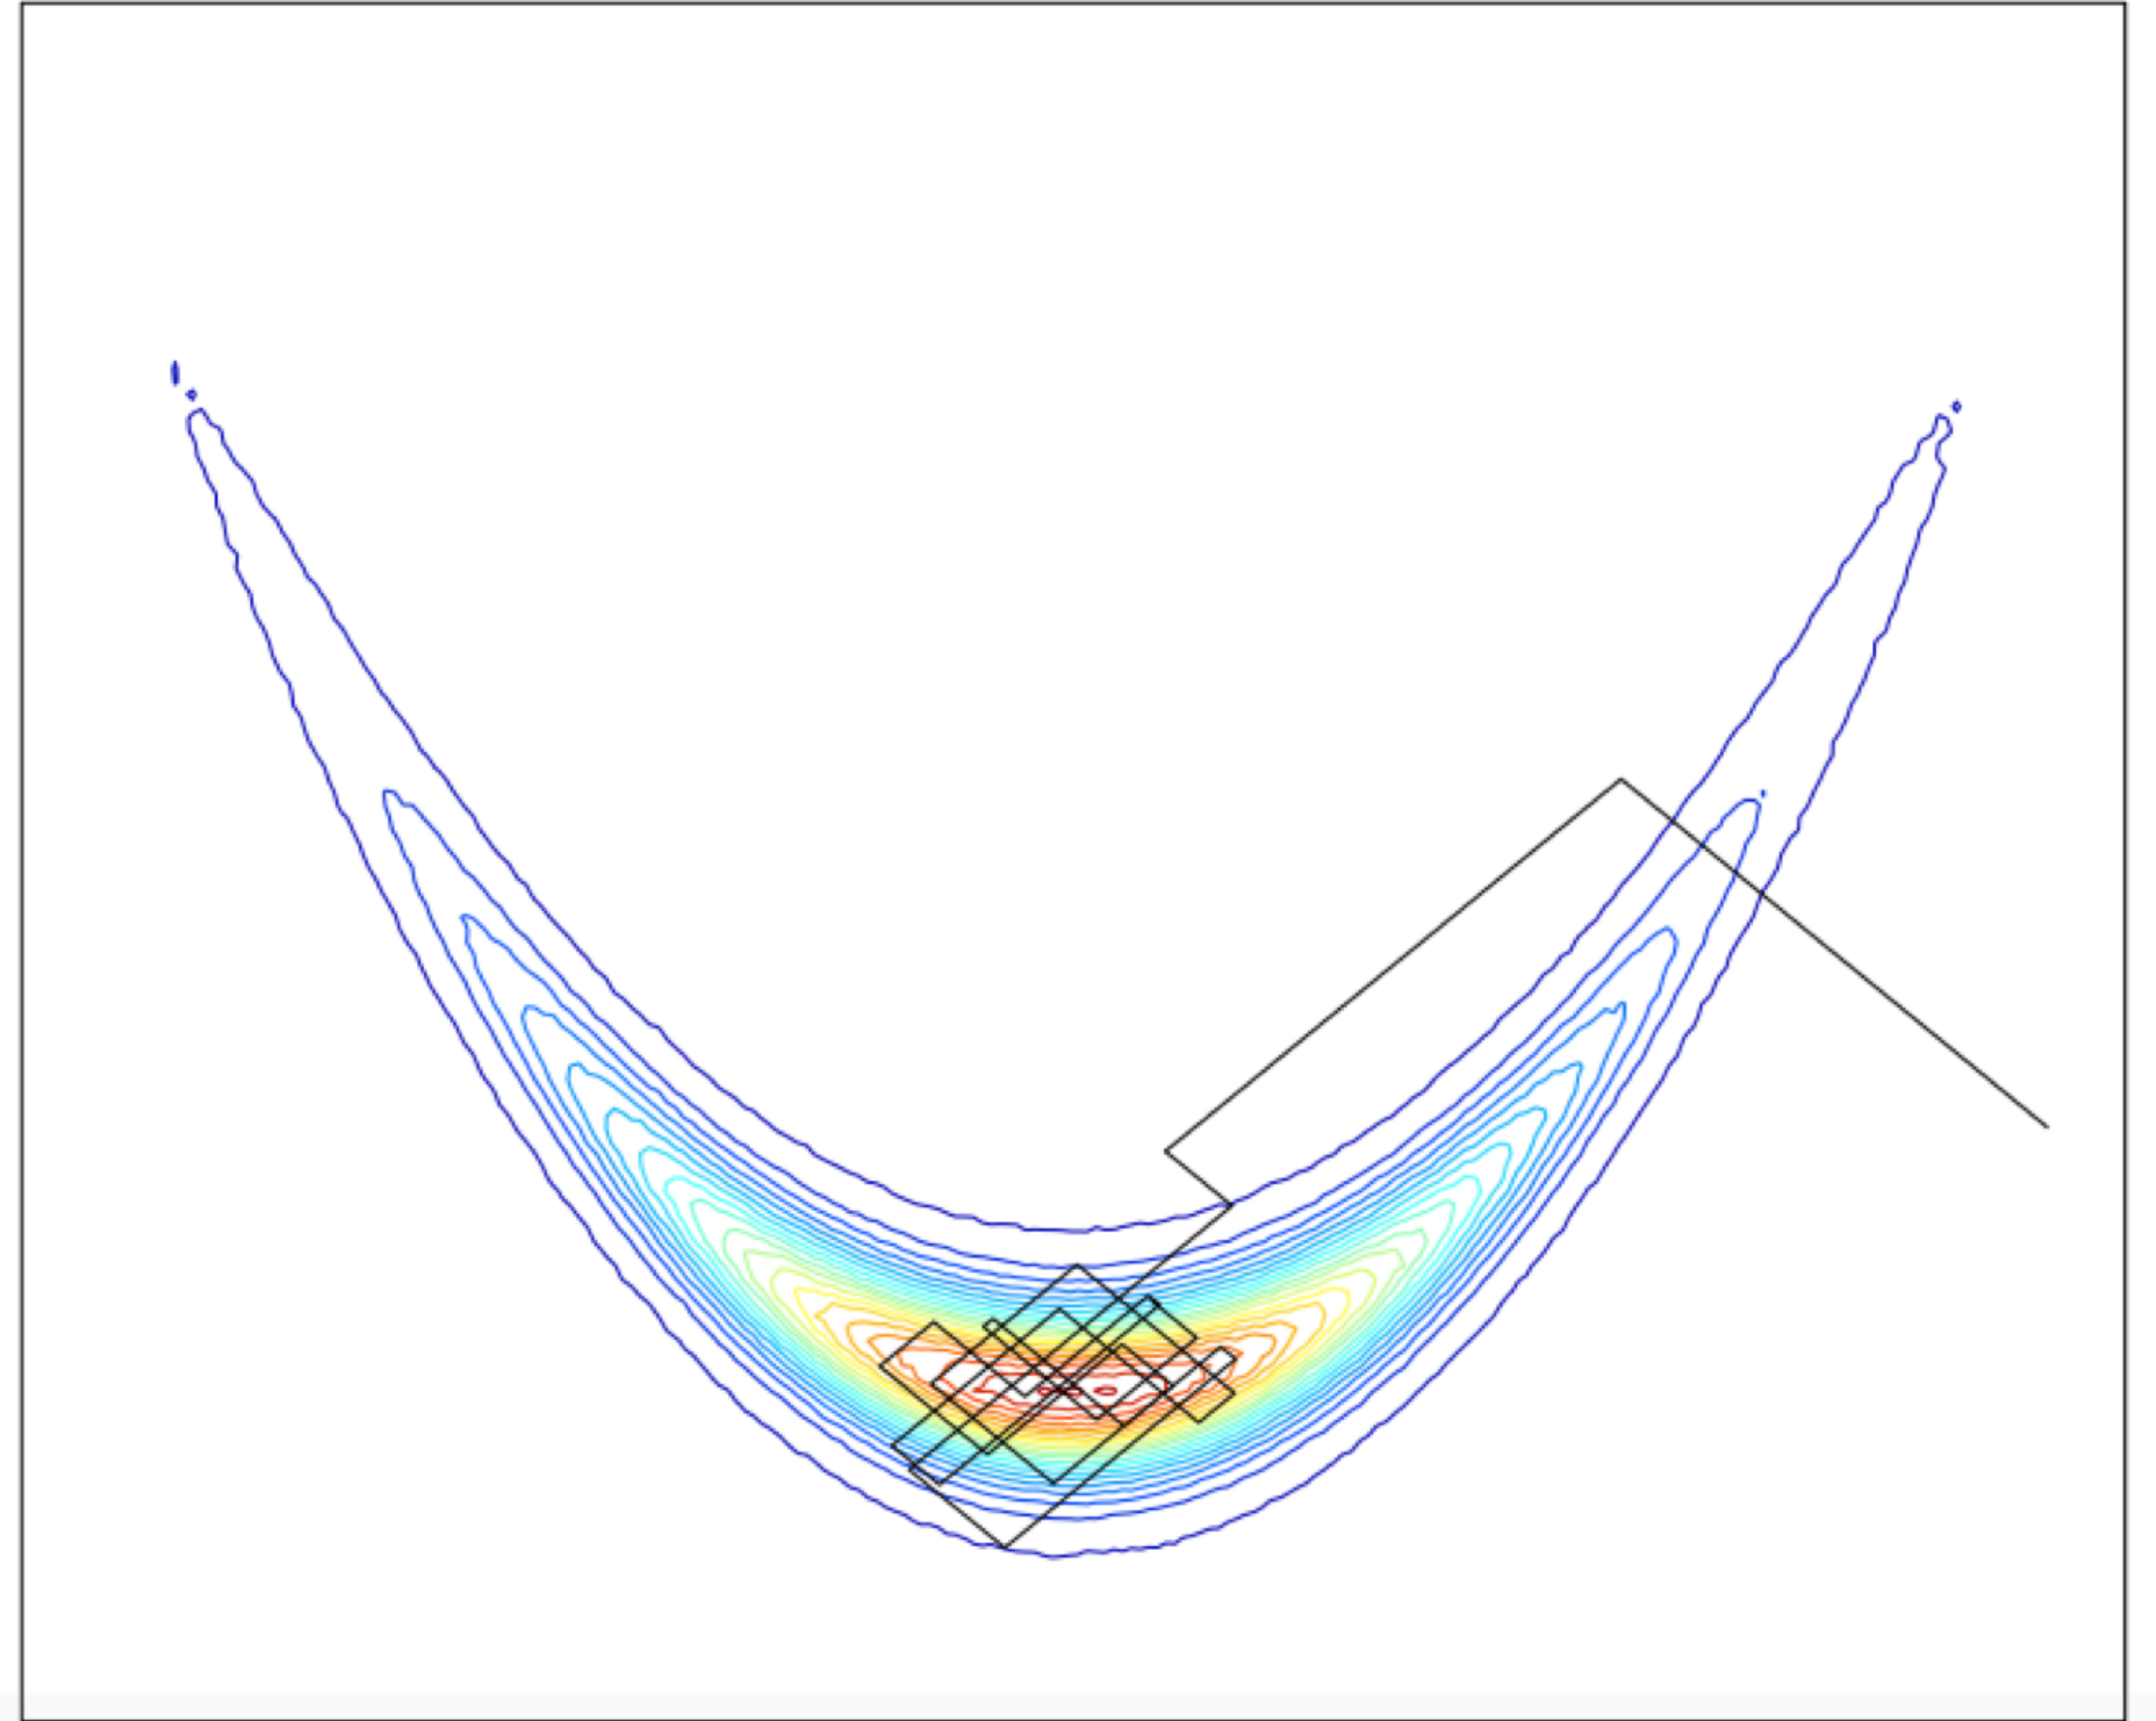
ZigZag Trajectory



Primer on PDMPs

- ‘Piecewise-Deterministic Markov Processes’ need


ZigZag Trajectory



Primer on PDMPs

- ‘Piecewise-Deterministic Markov Processes’ need
 - a vector field


ZigZag Trajectory



Primer on PDMPs

- ‘Piecewise-Deterministic Markov Processes’ need
 - a vector field
 - an ‘event rate’


ZigZag Trajectory



Primer on PDMPs

- ‘Piecewise-Deterministic Markov Processes’ need
 - a vector field
 - an ‘event rate’
 - a jump kernel


ZigZag Trajectory



Primer on PDMPs

- ‘Piecewise-Deterministic Markov Processes’ need
 - a vector field
 - an ‘event rate’
 - a jump kernel
- and then


ZigZag Trajectory



Primer on PDMPs

- ‘Piecewise-Deterministic Markov Processes’ need
 - a vector field
 - an ‘event rate’
 - a jump kernel
- and then
 - follow vector field


ZigZag Trajectory



Primer on PDMPs

- ‘Piecewise-Deterministic Markov Processes’ need
 - a vector field
 - an ‘event rate’
 - a jump kernel
- and then
 - follow vector field
 - experience event


ZigZag Trajectory



Primer on PDMPs

- ‘Piecewise-Deterministic Markov Processes’ need
 - a vector field
 - an ‘event rate’
 - a jump kernel
- and then
 - follow vector field
 - experience event
 - perform jump

ZigZag Trajectory



Examples of PDMPs for Monte Carlo

(today: just free transport ($\dot{x} = v, \dot{v} = 0$), others exist)

Examples of PDMPs for Monte Carlo

(today: just free transport ($\dot{x} = v, \dot{v} = 0$), others exist)

1. ‘Bouncy Particle Sampler’

- events at rate $\lambda(x, v) = \max(0, \langle v, \nabla U(x) \rangle)$
- jumps: specular reflection of v along contours of U

Examples of PDMPs for Monte Carlo

(today: just free transport ($\dot{x} = v, \dot{v} = 0$), others exist)

2. 'Zig-Zag Process'

- events at rate $\lambda_i(x, v) = \max \left(0, \langle v_i, \nabla_{x_i} U(x) \rangle \right)$
- jumps: flip v_i

Examples of PDMPs for Monte Carlo

(today: just free transport ($\dot{x} = v, \dot{v} = 0$), others exist)

3. 'Local' Bouncy Particle Sampler

- assume structured target: $\pi(x) = \exp \left(- \sum_a U_a(x_{\partial a}) \right)$
- events at rate $\lambda_a(x, v) = \max \left(0, \langle v_{\partial a}, \nabla U_a(x_{\partial a}) \rangle \right)$
- jumps: specular reflection of $v_{\partial a}$ along contours of U_a

Design of PDMP Methods

Design of PDMP Methods

- general characterisation of invariant PDMPs: too broad!
 - c.f. all invariant jump processes

Design of PDMP Methods

- general characterisation of invariant PDMPs: too broad!
 - c.f. all invariant jump processes
- meaningful restriction: invariant **and** 'skew-reversible'

Design of PDMP Methods

- general characterisation of invariant PDMPs: too broad!
 - c.f. all invariant jump processes
- meaningful restriction: invariant **and** ‘skew-reversible’
- abstraction: ‘time-enriched’ PDMP
 - process is $(x, \tau) \in \mathbb{R}^d \times \{\pm 1\}$
 - dynamics: $\dot{x} = \tau \cdot b(x)$
 - $\tau \sim$ ‘direction of time’; at jumps, flip τ

Design of PDMP Methods

- general characterisation of invariant PDMPs: too broad!
 - c.f. all invariant jump processes
- meaningful restriction: invariant **and** ‘skew-reversible’
- abstraction: ‘time-enriched’ PDMP
 - process is $(x, \tau) \in \mathbb{R}^d \times \{\pm 1\}$
 - dynamics: $\dot{x} = \tau \cdot b(x)$
 - $\tau \sim$ ‘direction of time’; at jumps, flip τ
- skew-reversibility: backwards in time, process looks *almost* the same

Basic Invariance Result

Basic Invariance Result

- Define ‘natural event rate’ $r(x) = \langle b(x), \nabla U(x) \rangle - \operatorname{div} b(x)$

Basic Invariance Result

- Define ‘natural event rate’ $r(x) = \langle b(x), \nabla U(x) \rangle - \operatorname{div} b(x)$
- Set $\lambda(x, \tau) = \max(0, \tau \cdot r(x)) + \gamma(x)$

Basic Invariance Result

- Define ‘natural event rate’ $r(x) = \langle b(x), \nabla U(x) \rangle - \operatorname{div} b(x)$
- Set $\lambda(x, \tau) = \max(0, \tau \cdot r(x)) + \gamma(x)$
- Let $Q^\tau(x \rightarrow \cdot)$ be reversible w.r.t. $\pi(x) \cdot \lambda(x, \tau)$

Basic Invariance Result

- Define ‘natural event rate’ $r(x) = \langle b(x), \nabla U(x) \rangle - \operatorname{div} b(x)$
- Set $\lambda(x, \tau) = \max(0, \tau \cdot r(x)) + \gamma(x)$
- Let $Q^\tau(x \rightarrow \cdot)$ be reversible w.r.t. $\pi(x) \cdot \lambda(x, \tau)$
- **Theorem:** necessary and sufficient for π -skew-reversibility

Basic Invariance Result

- Define ‘natural event rate’ $r(x) = \langle b(x), \nabla U(x) \rangle - \operatorname{div} b(x)$
- Set $\lambda(x, \tau) = \max(0, \tau \cdot r(x)) + \gamma(x)$
- Let $Q^\tau(x \rightarrow \cdot)$ be reversible w.r.t. $\pi(x) \cdot \lambda(x, \tau)$
- **Theorem:** necessary and sufficient for π -skew-reversibility
 - invariance follows

Structured Enhancement

Structured Enhancement

- recall ZZ, LBPS: event types are tailored to coordinates, target structure

Structured Enhancement

- recall ZZ, LBPS: event types are tailored to coordinates, target structure
- generalisation: *split* time-enriched PDMP

- dynamics: $\dot{x} = \sum_{\ell} \tau_{\ell} b_{\ell}(x), \tau_{\ell} \in \{\pm 1\}$

- decompose $r(x, \tau) = \sum_j r_j(x, \tau)$

- each term antisymmetric under a certain ‘flip’ \mathcal{F}_j of τ variables
- at ‘ j -jump’, apply j^{th} flip operator

Structured Invariance Result

Structured Invariance Result

- Set $\lambda_j(x, \tau) = \max \left(0, r_j(x, \tau) \right) + \gamma_j(x, \tau)$
 - require γ_j invariant under $\tau \mapsto -\tau, \tau \mapsto \mathcal{F}_j \tau$

Structured Invariance Result

- Set $\lambda_j(x, \tau) = \max \left(0, r_j(x, \tau) \right) + \gamma_j(x, \tau)$
 - require γ_j invariant under $\tau \mapsto -\tau, \tau \mapsto \mathcal{F}_j \tau$
- Let $Q_j^\tau(x \rightarrow \cdot)$ be reversible w.r.t. $\pi(x) \cdot \lambda_j(x, \tau)$
 - require $Q_j^\tau = Q_j^{-\mathcal{F}_j \tau}$

Structured Invariance Result

- Set $\lambda_j(x, \tau) = \max \left(0, r_j(x, \tau) \right) + \gamma_j(x, \tau)$
 - require γ_j invariant under $\tau \mapsto -\tau, \tau \mapsto \mathcal{F}_j \tau$
- Let $Q_j^\tau(x \rightarrow \cdot)$ be reversible w.r.t. $\pi(x) \cdot \lambda_j(x, \tau)$
 - require $Q_j^\tau = Q_j^{-\mathcal{F}_j \tau}$
- **Theorem:** necessary and sufficient for π -skew-reversibility
 - invariance follows

Conclusions

Conclusions

- dynamical systems provide natural principles for sampling algorithms

Conclusions

- dynamical systems provide natural principles for sampling algorithms
 - drift, diffuse, jump!

Conclusions

- dynamical systems provide natural principles for sampling algorithms
 - drift, diffuse, jump!
- can characterise π -invariant piecewise-deterministic Markov processes

Conclusions

- dynamical systems provide natural principles for sampling algorithms
 - drift, diffuse, jump!
- can characterise π -invariant piecewise-deterministic Markov processes
 - structured { targets / dynamics / events } easier to analyse

Conclusions

- dynamical systems provide natural principles for sampling algorithms
 - drift, diffuse, jump!
- can characterise π -invariant piecewise-deterministic Markov processes
 - structured { targets / dynamics / events } easier to analyse
 - simplifies design of new methods

Conclusions

- dynamical systems provide natural principles for sampling algorithms
 - drift, diffuse, jump!
- can characterise π -invariant piecewise-deterministic Markov processes
 - structured { targets / dynamics / events } easier to analyse
 - simplifies design of new methods
- ongoing: simplifying implementation

thanks!

Design of Sampling Dynamics

(ODE, SDE case)

- bare minimum: dynamics should leave π invariant
 - linear condition on generator of process (simple to check, in principle)
- for pure ODE: specify skew-symmetric $Q(x)$, then define
 - $b(x) = Q(x) \nabla \log \pi(x) + \operatorname{div} Q(x)$
 - evolve by $\dot{x} = b(x)$
- for reversible SDE: specify PSD $D(x)$, then define
 - $b(x) = D(x) \nabla \log \pi(x) + \operatorname{div} D(x)$
 - evolve by $dx = b(x) dt + \sqrt{2D(x)} dW$

Design of Sampling Dynamics

(ODE, SDE case)

- general SDE = (pure ODE) + (reversible SDE)
 - so: pick $(Q(x), D(x))$, and you're good to go!
- result is ~ old:
 - versions known in stat phys since ~1970s
 - re-popularised by Ma-Chen-Wu-Fox in ML (stochastic gradients)
 - expanded by Barp-Betancourt-Takao-Arnaudon-Girolami (geometry)
- but ... PDMPs?