

PDMPs with ODE Dynamics

Sam Power
(joint with Sergio Bacallado)



Algorithms & Computationally Intensive Inference Seminar
University of Warwick

sp825@cam.ac.uk

June 6, 2019

- 1 PDMPs
- 2 PDMPs for MCMC
- 3 Construction of Algorithms
- 4 Remarks, Open Questions, Takeaways

- Informally: Deterministic dynamics + Jump Process

- Informally: Deterministic dynamics + Jump Process
- Stochastic process Z_t which


- Informally: Deterministic dynamics + Jump Process
- Stochastic process Z_t which
 - 1 Follows a *deterministic* path, until

- Informally: Deterministic dynamics + Jump Process
- Stochastic process Z_t which
 - 1 Follows a *deterministic* path, until
 - 2 An *event* occurs, at a certain rate, upon which

- Informally: Deterministic dynamics + Jump Process
- Stochastic process Z_t which
 - 1 Follows a *deterministic* path, until
 - 2 An *event* occurs, at a certain rate, upon which
 - 3 The position *jumps*, and then

- Informally: Deterministic dynamics + Jump Process
- Stochastic process Z_t which
 - 1 Follows a *deterministic* path, until
 - 2 An *event* occurs, at a certain rate, upon which
 - 3 The position *jumps*, and then
 - 4 Resumes following the deterministic path

- Informally: Deterministic dynamics + Jump Process
- Stochastic process Z_t which
 - 1 Follows a *deterministic* path, until
 - 2 An *event* occurs, at a certain rate, upon which
 - 3 The position *jumps*, and then
 - 4 Resumes following the deterministic path



Specifying a PDMP

- Today: PDMPs from ODEs

Specifying a PDMP

- Today: PDMPs from ODEs
 - Vector field $\phi(z)$
 - Use dynamics $\frac{dz}{dt} = \phi(z)$

- Today: PDMPs from ODEs
 - Vector field $\phi(z)$
 - Use dynamics $\frac{dz}{dt} = \phi(z)$
 - Event rate $\lambda(z) \geq 0$
 - Dictates how often events happen (inhomogeneous Poisson process)

- Today: PDMPs from ODEs
 - Vector field $\phi(z)$
 - Use dynamics $\frac{dz}{dt} = \phi(z)$
 - Event rate $\lambda(z) \geq 0$
 - Dictates how often events happen (inhomogeneous Poisson process)
 - Transition dynamics $Q(z \rightarrow dz')$
 - Dictates what happens at events (Markov jump kernel)

- Want $\pi(dx)$, but work on extended target:

- Want $\pi(dx)$, but work on extended target:
 - Set $z = (x, v)$.

- Want $\pi(dx)$, but work on extended target:
 - Set $z = (x, v)$.
 - Choose your own $\psi(dv)$.

- Want $\pi(dx)$, but work on extended target:
 - Set $z = (x, v)$.
 - Choose your own $\psi(dv)$.
 - Target is then $\mu(dz) = \pi(dx)\psi(dv)$.

- Want $\pi(dx)$, but work on extended target:
 - Set $z = (x, v)$.
 - Choose your own $\psi(dv)$.
 - Target is then $\mu(dz) = \pi(dx)\psi(dv)$.
- Typically, jumps fix $x \rightsquigarrow X_t$ has continuous sample paths.

- Want $\pi(dx)$, but work on extended target:
 - Set $z = (x, v)$.
 - Choose your own $\psi(dv)$.
 - Target is then $\mu(dz) = \pi(dx)\psi(dv)$.
- Typically, jumps fix $x \rightsquigarrow X_t$ has continuous sample paths.
- **Question:**

Given target measure μ , vector field ϕ , (1)

how can I build (λ, Q) to sample μ ? (2)

Examples of PDMPs

- Bouncy Particle Sampler

Examples of PDMPs

- Bouncy Particle Sampler

- $\dot{x} = v, \dot{v} = 0$

- Bouncy Particle Sampler
 - $\dot{x} = v, \dot{v} = 0$
 - $\lambda(x, v) = \langle v, -\nabla_x \log \pi(x) \rangle_+$

- Bouncy Particle Sampler
 - $\dot{x} = v, \dot{v} = 0$
 - $\lambda(x, v) = \langle v, -\nabla_x \log \pi(x) \rangle_+$
 - Jump by reflecting velocity

- Bouncy Particle Sampler
 - $\dot{x} = v, \dot{v} = 0$
 - $\lambda(x, v) = \langle v, -\nabla_x \log \pi(x) \rangle_+$
 - Jump by reflecting velocity
- Zig Zag Process

- Bouncy Particle Sampler
 - $\dot{x} = v, \dot{v} = 0$
 - $\lambda(x, v) = \langle v, -\nabla_x \log \pi(x) \rangle_+$
 - Jump by reflecting velocity
- Zig Zag Process
 - $\dot{x} = v, \dot{v} = 0$

- Bouncy Particle Sampler

- $\dot{x} = v, \dot{v} = 0$
- $\lambda(x, v) = \langle v, -\nabla_x \log \pi(x) \rangle_+$
- Jump by reflecting velocity

- Zig Zag Process

- $\dot{x} = v, \dot{v} = 0$
- $\lambda_i(x, v) = \langle v_i, -\nabla_{x_i} \log \pi(x) \rangle_+$

- Bouncy Particle Sampler

- $\dot{x} = v, \dot{v} = 0$
- $\lambda(x, v) = \langle v, -\nabla_x \log \pi(x) \rangle_+$
- Jump by reflecting velocity

- Zig Zag Process

- $\dot{x} = v, \dot{v} = 0$
- $\lambda_i(x, v) = \langle v_i, -\nabla_{x_i} \log \pi(x) \rangle_+$
- Jump by flipping i^{th} velocity

- Reversibility
 - Much MCMC work built on reversible methods
 - PDMPs are generally non-reversible
 - To design algorithms, *locality* is the important part

- Reversibility
 - Much MCMC work built on reversible methods
 - PDMPs are generally non-reversible
 - To design algorithms, *locality* is the important part
- Symmetry
 - Existing PDMPs are highly symmetric (BPS, ZZ)
 - A priori, not necessary to have symmetry
 - Want to be able to use *all* ODEs!

Time-Augmented PDMPs

- Idea:

Time-Augmented PDMPs

- Idea:

- 1 Introduce 'direction of time' variable $\tau \in \{\pm 1\}$

Time-Augmented PDMPs

- Idea:

- 1 Introduce 'direction of time' variable $\tau \in \{\pm 1\}$
- 2 Target $\tilde{\mu}(dz, d\tau) = \mu(dz)R(d\tau)$.

Time-Augmented PDMPs

- Idea:

- ① Introduce 'direction of time' variable $\tau \in \{\pm 1\}$

- ② Target $\tilde{\mu}(dz, d\tau) = \mu(dz)R(d\tau)$.

- Write $\phi(z, \tau) = \tau \cdot \phi(z)$; use dynamics $\frac{dz}{dt} = \phi(z, \tau)$

Time-Augmented PDMPs

- Idea:
 - ① Introduce 'direction of time' variable $\tau \in \{\pm 1\}$
 - ② Target $\tilde{\mu}(dz, d\tau) = \mu(dz)R(d\tau)$.
- Write $\phi(z, \tau) = \tau \cdot \phi(z)$; use dynamics $\frac{dz}{dt} = \phi(z, \tau)$
 - Solve system forwards **and** backwards in time

Time-Augmented PDMPs

- Idea:
 - ① Introduce 'direction of time' variable $\tau \in \{\pm 1\}$
 - ② Target $\tilde{\mu}(dz, d\tau) = \mu(dz)R(d\tau)$.
- Write $\phi(z, \tau) = \tau \cdot \phi(z)$; use dynamics $\frac{dz}{dt} = \phi(z, \tau)$
 - Solve system forwards **and** backwards in time
- Let $\lambda = \lambda(z, \tau)$

Time-Augmented PDMPs

- Idea:
 - ① Introduce 'direction of time' variable $\tau \in \{\pm 1\}$
 - ② Target $\tilde{\mu}(dz, d\tau) = \mu(dz)R(d\tau)$.
- Write $\phi(z, \tau) = \tau \cdot \phi(z)$; use dynamics $\frac{dz}{dt} = \phi(z, \tau)$
 - Solve system forwards **and** backwards in time
- Let $\lambda = \lambda(z, \tau)$
- Stipulate that, at events, $\tau \mapsto -\tau$, i.e.

$$Q((z, \tau) \rightarrow (dz', d\tau')) = Q^\tau(z \rightarrow dz') \cdot \delta(-\tau, d\tau') \quad (3)$$

Time-Augmented PDMPs

- Idea:
 - ① Introduce 'direction of time' variable $\tau \in \{\pm 1\}$
 - ② Target $\tilde{\mu}(dz, d\tau) = \mu(dz)R(d\tau)$.
- Write $\phi(z, \tau) = \tau \cdot \phi(z)$; use dynamics $\frac{dz}{dt} = \phi(z, \tau)$
 - Solve system forwards **and** backwards in time
- Let $\lambda = \lambda(z, \tau)$
- Stipulate that, at events, $\tau \mapsto -\tau$, i.e.

$$Q((z, \tau) \rightarrow (dz', d\tau')) = Q^\tau(z \rightarrow dz') \cdot \delta(-\tau, d\tau') \quad (3)$$

- 'Trajectorial Reversibility' \leadsto checking exactness becomes *local*!

Time-Augmented PDMPs

- Idea:
 - ① Introduce 'direction of time' variable $\tau \in \{\pm 1\}$
 - ② Target $\tilde{\mu}(dz, d\tau) = \mu(dz)R(d\tau)$.
- Write $\phi(z, \tau) = \tau \cdot \phi(z)$; use dynamics $\frac{dz}{dt} = \phi(z, \tau)$
 - Solve system forwards **and** backwards in time
- Let $\lambda = \lambda(z, \tau)$
- Stipulate that, at events, $\tau \mapsto -\tau$, i.e.

$$Q((z, \tau) \rightarrow (dz', d\tau')) = Q^\tau(z \rightarrow dz') \cdot \delta(-\tau, d\tau') \quad (3)$$

- 'Trajectorial Reversibility' \leadsto checking exactness becomes *local*!
 - 'in at z forwards in time = out at z backwards in time'

Choice of Event Rate (1)

- Consider 'probability current'

$$r(z, \tau) \triangleq \underbrace{\langle \nabla H(z), \phi(z, \tau) \rangle}_{\text{Energy Gain}} - \underbrace{\text{div}_z \phi(z, \tau)}_{\text{Compressibility Penalty}} \quad (4)$$

Choice of Event Rate (1)

- Consider 'probability current'

$$r(z, \tau) \triangleq \underbrace{\langle \nabla H(z), \phi(z, \tau) \rangle}_{\text{Energy Gain}} - \underbrace{\text{div}_z \phi(z, \tau)}_{\text{Compressibility Penalty}} \quad (4)$$

- Define 'natural' event rate as

$$\lambda^0(z, \tau) = (r(z, \tau))_+ \quad (5)$$

where $(u)_+ = \max(0, u)$

Choice of Event Rate (1)

- Consider 'probability current'

$$r(z, \tau) \triangleq \underbrace{\langle \nabla H(z), \phi(z, \tau) \rangle}_{\text{Energy Gain}} - \underbrace{\text{div}_z \phi(z, \tau)}_{\text{Compressibility Penalty}} \quad (4)$$

- Define 'natural' event rate as

$$\lambda^0(z, \tau) = (r(z, \tau))_+ \quad (5)$$

where $(u)_+ = \max(0, u)$

- Let $\gamma(z) \geq 0$ be some 'refreshment rate'.

Choice of Event Rate (1)

- Consider 'probability current'

$$r(z, \tau) \triangleq \underbrace{\langle \nabla H(z), \phi(z, \tau) \rangle}_{\text{Energy Gain}} - \underbrace{\text{div}_z \phi(z, \tau)}_{\text{Compressibility Penalty}} \quad (4)$$

- Define 'natural' event rate as

$$\lambda^0(z, \tau) = (r(z, \tau))_+ \quad (5)$$

where $(u)_+ = \max(0, u)$

- Let $\gamma(z) \geq 0$ be some 'refreshment rate'.
- We will take $\lambda(z, \tau) = \lambda^0(z, \tau) + \gamma(z)$

- Define 'jump measure':

$$J^\tau(dz) \propto \mu(dz)\lambda(z, \tau) \quad (6)$$

- Define 'jump measure':

$$J^\tau(dz) \propto \mu(dz)\lambda(z, \tau) \quad (6)$$

- Want trajectorial reversibility

- Define 'jump measure':

$$J^\tau(dz) \propto \mu(dz)\lambda(z, \tau) \quad (6)$$

- Want trajectorial reversibility
 - \implies Need jump chain reversible w.r.t. jump measure

- Define 'jump measure':

$$J^\tau(dz) \propto \mu(dz)\lambda(z, \tau) \quad (6)$$

- Want trajectorial reversibility
 - \implies Need jump chain reversible w.r.t. jump measure
 - \leadsto Choose $q^\tau(z \rightarrow dz')$ to be J^τ -reversible

Putting together the ingredients

Theorem

If (ϕ, λ, Q) are chosen in this way, then the resulting PDMP is trajectoryally reversible, and admits $\tilde{\mu}$ as a stationary measure.

Putting together the ingredients

Theorem

If (ϕ, λ, Q) are chosen in this way, then the resulting PDMP is trajectoryally reversible, and admits $\tilde{\mu}$ as a stationary measure.

Theorem

If (ϕ, λ, Q) is a trajectoryally-reversible, $\tilde{\mu}$ -stationary TA-PDMP, then $\exists \gamma \geq 0$ such that

$$\lambda(z, \tau) = \lambda^0(z, \tau) + \gamma(z) \quad (7)$$

and for $\tau \in \{\pm 1\}$, Q^τ is J^τ -reversible

Putting together the ingredients

Theorem

If (ϕ, λ, Q) are chosen in this way, then the resulting PDMP is trajectorially reversible, and admits $\tilde{\mu}$ as a stationary measure.

Theorem

If (ϕ, λ, Q) is a trajectorially-reversible, $\tilde{\mu}$ -stationary TA-PDMP, then $\exists \gamma \geq 0$ such that

$$\lambda(z, \tau) = \lambda^0(z, \tau) + \gamma(z) \quad (7)$$

and for $\tau \in \{\pm 1\}$, Q^τ is J^τ -reversible

- Comments on proof

Split PDMPs (1)

- Many PDMPs in use have different types of event

Split PDMPs (1)

- Many PDMPs in use have different types of event
 - Refreshment
 - Zig-Zag
 - Local BPS (Factor Graph)
 - Subsampling
 - ...

Split PDMPs (1)

- Many PDMPs in use have different types of event
 - Refreshment
 - Zig-Zag
 - Local BPS (Factor Graph)
 - Subsampling
 - ...
- Each event type affects different parts of the system

Split PDMPs (1)

- Many PDMPs in use have different types of event
 - Refreshment
 - Zig-Zag
 - Local BPS (Factor Graph)
 - Subsampling
 - ...
- Each event type affects different parts of the system
- Key point: Different event types correspond to *decompositions* of r

Split PDMPs (2)

- $z = (z_1, \dots, z_D), \tau = (\tau_1, \dots, \tau_D) \in \{\pm 1\}^D$
- $\phi(z, \tau) = \tau \odot \phi(z) = (\tau_1 \phi_1(z), \dots, \tau_D \phi_D(z))$

Split PDMPs (2)

- $z = (z_1, \dots, z_D)$, $\tau = (\tau_1, \dots, \tau_D) \in \{\pm 1\}^D$
- $\phi(z, \tau) = \tau \odot \phi(z) = (\tau_1 \phi_1(z), \dots, \tau_D \phi_D(z))$
- Assume decomposition

$$r(z, \tau) = \sum_{j=1}^M r_j(z, \tau) \quad (8)$$

and existence of involutions $\mathcal{F}_j : \{\pm 1\}^D \rightarrow \{\pm 1\}^D$ such that

$$r_j(z, \mathcal{F}_j(\tau)) = -r_j(z, \tau) \quad (9)$$

Split PDMPs (2)

- $z = (z_1, \dots, z_D)$, $\tau = (\tau_1, \dots, \tau_D) \in \{\pm 1\}^D$
- $\phi(z, \tau) = \tau \odot \phi(z) = (\tau_1 \phi_1(z), \dots, \tau_D \phi_D(z))$
- Assume decomposition

$$r(z, \tau) = \sum_{j=1}^M r_j(z, \tau) \quad (8)$$

and existence of involutions $\mathcal{F}_j : \{\pm 1\}^D \rightarrow \{\pm 1\}^D$ such that

$$r_j(z, \mathcal{F}_j(\tau)) = -r_j(z, \tau) \quad (9)$$

- Events of type j happen at rate $\lambda_j(z, \tau)$
 - and then jump according to $Q_j^\tau(z \rightarrow dz') \cdot \delta(\mathcal{F}_j(\tau), d\tau')$

Making Split-PDMPs work (1)

- Define

$$\lambda_j^0(z, \tau) = (r_j(z, \tau))_+ \quad (10)$$

$$\lambda_j(z, \tau) = \lambda_j^0(z, \tau) + \gamma_j(z, \tau) \quad (11)$$

Making Split-PDMPs work (1)

- Define

$$\lambda_j^0(z, \tau) = (r_j(z, \tau))_+ \quad (10)$$

$$\lambda_j(z, \tau) = \lambda_j^0(z, \tau) + \gamma_j(z, \tau) \quad (11)$$

- Define

$$J_j^\tau(dz) \propto \mu(dz) \lambda_j(z, \tau) \quad (12)$$

and for each $\tau \in \{\pm 1\}^D$, take Q_j^τ to be J_j^τ -reversible.

Making Split-PDMPs work (1)

- Define

$$\lambda_j^0(z, \tau) = (r_j(z, \tau))_+ \quad (10)$$

$$\lambda_j(z, \tau) = \lambda_j^0(z, \tau) + \gamma_j(z, \tau) \quad (11)$$

- Define

$$J_j^\tau(dz) \propto \mu(dz) \lambda_j(z, \tau) \quad (12)$$

and for each $\tau \in \{\pm 1\}^D$, take Q_j^τ to be J_j^τ -reversible.

Theorem

This leads to trajectorially-reversible, $\tilde{\mu}$ -stationary Split PDMPs.

Making Split-PDMPs work (1)

- Define

$$\lambda_j^0(z, \tau) = (r_j(z, \tau))_+ \quad (10)$$

$$\lambda_j(z, \tau) = \lambda_j^0(z, \tau) + \gamma_j(z, \tau) \quad (11)$$

- Define

$$J_j^\tau(dz) \propto \mu(dz) \lambda_j(z, \tau) \quad (12)$$

and for each $\tau \in \{\pm 1\}^D$, take Q_j^τ to be J_j^τ -reversible.

Theorem

This leads to trajectoryally-reversible, $\tilde{\mu}$ -stationary Split PDMPs.

Theorem

Given a fixed splitting, all trajectoryally-reversible, $\tilde{\mu}$ -stationary Split PDMPs take this form.

Algorithm Design Pipeline (1)

- Non-negotiable: we want samples from $\pi(dx)$.

Algorithm Design Pipeline (1)

- Non-negotiable: we want samples from $\pi(dx)$.
 - 1 Decide on v .
 - 2 Decide on ϕ .
 - 3 Decide on $\psi(dv)$ (and hence μ).

Algorithm Design Pipeline (1)

- Non-negotiable: we want samples from $\pi(dx)$.
 - 1 Decide on v .
 - 2 Decide on ϕ .
 - 3 Decide on $\psi(dv)$ (and hence μ).
 - 4 Write down r , decide on a splitting.

Algorithm Design Pipeline (1)

- Non-negotiable: we want samples from $\pi(dx)$.
 - ➊ Decide on v .
 - ➋ Decide on ϕ .
 - ➌ Decide on $\psi(dv)$ (and hence μ).
 - ➍ Write down r , decide on a splitting.
 - ➎ Write down λ^0 , decide on γ (and hence λ).

Algorithm Design Pipeline (1)

- Non-negotiable: we want samples from $\pi(dx)$.
 - ➊ Decide on v .
 - ➋ Decide on ϕ .
 - ➌ Decide on $\psi(dv)$ (and hence μ).
 - ➍ Write down r , decide on a splitting.
 - ➎ Write down λ^0 , decide on γ (and hence λ).
 - ➏ Decide on Q .

Algorithm Design Pipeline (2)

- Choosing Q is often least obvious; order of preference:

Algorithm Design Pipeline (2)

- Choosing Q is often least obvious; order of preference:
 - 1 Sample from J^τ directly.

Algorithm Design Pipeline (2)

- Choosing Q is often least obvious; order of preference:
 - 1 Sample from J^τ directly.
 - 2 Sample from its restriction to a finite set. (e.g. BPS)

- Choosing Q is often least obvious; order of preference:
 - 1 Sample from J^τ directly.
 - 2 Sample from its restriction to a finite set. (e.g. BPS)
 - 3 (Use a Metropolis-Hastings step).

Algorithm Design Pipeline (2)

- Choosing Q is often least obvious; order of preference:
 - 1 Sample from J^τ directly.
 - 2 Sample from its restriction to a finite set. (e.g. BPS)
 - 3 (Use a Metropolis-Hastings step).
- Choosing ψ could make a big difference; dictates μ .
 - Can have $\psi(dv|x)$ (relatively unexplored)

Algorithm Design Pipeline (2)

- Choosing Q is often least obvious; order of preference:
 - 1 Sample from J^τ directly.
 - 2 Sample from its restriction to a finite set. (e.g. BPS)
 - 3 (Use a Metropolis-Hastings step).
- Choosing ψ could make a big difference; dictates μ .
 - Can have $\psi(dv|x)$ (relatively unexplored)
- Choosing ϕ : some room for creativity here.

- Andrieu, Livingstone (2018): Peskun-type ordering for (some) PDMPs

- Andrieu, Livingstone (2018): Peskun-type ordering for (some) PDMPs
 - **Conjecture:** Split as little as possible

- Andrieu, Livingstone (2018): Peskun-type ordering for (some) PDMPs
 - **Conjecture:** Split as little as possible
 - **Conjecture:** Refresh as little as possible

- Andrieu, Livingstone (2018): Peskun-type ordering for (some) PDMPs
 - **Conjecture:** Split as little as possible
 - **Conjecture:** Refresh as little as possible
 - Pinch of salt / 'Pre-Asymptopia': Maire, Vialaret (2018)

- Andrieu, Livingstone (2018): Peskun-type ordering for (some) PDMPs
 - **Conjecture:** Split as little as possible
 - **Conjecture:** Refresh as little as possible
 - Pinch of salt / 'Pre-Asymptopia': Maire, Vialaret (2018)
- Implementation remains challenging
 - Splittings may help

- Andrieu, Livingstone (2018): Peskun-type ordering for (some) PDMPs
 - **Conjecture:** Split as little as possible
 - **Conjecture:** Refresh as little as possible
 - Pinch of salt / 'Pre-Asymptopia': Maire, Vialaret (2018)
- Implementation remains challenging
 - Splittings may help
- *Speculation:* Better dynamics $\phi \rightsquigarrow$ opportunities

- Andrieu, Livingstone (2018): Peskun-type ordering for (some) PDMPs
 - **Conjecture:** Split as little as possible
 - **Conjecture:** Refresh as little as possible
 - Pinch of salt / 'Pre-Asymptopia': Maire, Vialaret (2018)
- Implementation remains challenging
 - Splittings may help
- *Speculation:* Better dynamics $\phi \rightsquigarrow$ opportunities
- *Curiosity:* Other types of augmentation?

Thank you!