

Explicit convergence bounds for Metropolis Markov chains

Isoperimetry, Spectral Gaps and Profiles

Sam Power

University of Bristol

27 January, 2023

Links & Acknowledgements

- ✿ Main paper today: arXiv 2211.08959;
- ✿ Related: arXiv 2208.05239
- ✿ All joint work with
 - ▶ Christophe Andrieu (Bristol)
 - ▶ Anthony Lee (Bristol)
 - ▶ Andi Q. Wang (Bristol \rightsquigarrow Warwick)
- ✿ Funded by Bayes4Health EPSRC Grant

Setting: Task

- ✂ Motivating task: understanding structured probability distributions in high-dimensional spaces
 - ▶ posterior inference in Bayesian statistics
 - ▶ latent variable models, hidden Markov models
 - ▶ generative modeling
 - ▶ non-convex optimisation
 - ▶ ...

Markov Chain Monte Carlo (MCMC)

- ✂ Task: Generate approximate samples from a probability distribution π to which we have *limited access*.
- ✂ MCMC: An iterative approach to this task.
 - ▶ Simulate a time-homogeneous Markov chain $(X_n)_{n \geq 0}$ such that

$$\text{Law}(X_n) \rightarrow \pi \text{ as } n \rightarrow \infty.$$

(and hopefully, quickly)

- ✂ Current status:
 - ▶ Mature algorithmic field, many ‘correct’ solutions are known and practical.
 - ▶ Quantitative convergence theory is *challenging; important*.
 - ▶ ‘Is (this algorithm) { performant, reliable, preferable, ... } ?’
 - ▶ ‘Given π , which algorithm do I choose?’

Random Walk Metropolis

✂ Today: Study the *Random Walk Metropolis* (RWM) algorithm

- ▶ Only requires access to density of π , up to a multiplicative constant (typical).
- ▶ Widely-used, simple, 'representative'

1. At x ,

1.1 Propose $x' \sim \mathcal{N}(x, \sigma^2 \cdot I_d)$.

1.2 Evaluate $r(x, x') = \frac{\pi(x')}{\pi(x)}$.

1.3 With probability $\min\{1, r(x, x')\}$, move to x' ; otherwise, remain at x .

✂ Leaves π invariant, ergodic under mild conditions, exponentially so under tail conditions.

Some Notation

✿ $Q(x, dx') = \mathcal{N}(dx'; x, \sigma^2 \cdot I_d).$

✿ $\alpha(x, x') := \min \{1, r(x, x')\}.$

✿ $\alpha(x) = \int Q(x, dx') \alpha(x, x').$

✿ The RWM kernel P is given by

$$P(x, dx') = Q(x, dx') \cdot \alpha(x, x') \\ + (1 - \alpha(x)) \cdot \delta(x, dx')$$

Convergence Analysis of RWM

- ✿ ‘Soft’ analysis: Exponential convergence \Leftrightarrow Lighter-than-Exponential Tails.
- ✿ ‘Optimal scaling’ analysis: control acceptance rate to optimise efficiency.
- ✿ ‘Modern’ analysis: log-concavity of target, ‘optimisation-style’ proofs
 - ▶ (and variations on this)
- ✿ Today: synthesis of the above.

Some Comments on Results

- ✂ Despite ubiquity, sharp complexity analysis of RWM has long been open.
 - ▶ Better than rejection sampling, quadrature, \dots ?
- ✂ We obtain a convincing complexity analysis with
 - ▶ sharp dependence on the dimension of the problem
 - ▶ conjecturally sharp dependence on the conditioning of the problem
- ✂ Our proof techniques are remarkably robust, and largely new to this area
- ✂ Gives a relatively complete resolution to the question of RWM's mixing

Main Results

✂ Suppose that

- ▶ Target is $\pi(x) \propto \exp(-U(x))$,
- ▶ U is m -strongly convex, L -smooth,
- ▶ Write $\kappa = L/m$ ('condition number').

✂ Run RWM with $\sigma = v \cdot (L \cdot d)^{-1/2}$.

✂ Then,

1. Acceptance rate satisfies $\alpha(x) \geq \alpha_0 := \frac{1}{2} \cdot \exp(-\frac{1}{2}v^2)$.
2. Spectral gap satisfies $\gamma_P \geq c(v) \cdot (\kappa \cdot d)^{-1}$.
3. L^2 mixing time satisfies $T_*(\varepsilon) \lesssim \kappa \cdot d \cdot \log\left(\frac{\kappa \cdot d}{\varepsilon}\right)$

✂ Paper contains tools which imply simple bounds for much wider class of targets.

✂ Today: demystify those tools.

Proof Overview

✿ Roughly:

1. Large-Scale Properties of Target
2. + Small-Scale Properties of Sampler
3. \rightsquigarrow Good Mixing.

✿ Precisely:

- ▶ ‘Isoperimetric’ Profile of Target
- ▶ + ‘Close Coupling’ of Kernels
- ▶ \rightsquigarrow Isoperimetric Profile of *Markov Chain*
 - ▶ \rightsquigarrow Good Mixing (in L^2).

✿ True for fairly general Markov chains on metric spaces.

✿ For RWM *in particular*:

- ▶ ‘Metropolis-type’ + Acceptance Control \rightsquigarrow Close Coupling.

✿ I will explain all of these terms.

Isoperimetry, Conductance, and Escapes

- ✿ For 'local' Markov chains, a powerful tool of analysis is the 'conductance' method.
- ✿ The core idea is that if a chain cannot get stuck badly in a set of small mass, then the chain must be mixing well.
- ✿ Quantitatively: for any (small) set A , [the flow of the chain out of A and into A^c] is comparable to [the mass of A].
- ✿ If this condition holds, then the chain is mixing well.
 - ▶ Under some conditions, this is a theorem.
 - ▶ Weaker and stronger versions of this property are also useful.
 - ▶ These each lead to their own theorems.

Conductance Methods for Markov Chains

✿ Consider for $A \subseteq \mathbb{R}^d$

$$\pi(A) := \int_{x \in A} \pi(x) \, dx$$
$$\pi \otimes P(A \times A^c) := \int_{x \in A, y \in A^c} \pi(x) P(x, y) \, dx dy.$$

- ✿ If $\pi \otimes P(A \times A^c) \geq c \cdot \pi(A)$, then $P(X_1 \notin A \mid X_0 \in A) \geq c$,
 - ▶ so if $c \gg 0$, then the set A is easy for P to escape.
- ✿ If every set A is easy for P to escape, then P cannot get stuck ...
 - ▶ ... and hence must converge quickly.

Isoperimetric Profiles of Markov Chains

✿ Define

$$I_{\pi,P}(p) := \inf \left\{ \pi \otimes P \left(A \times A^c \right) : \pi(A) = p \right\}$$

✿ ‘How hard is it for *this chain* to leave sets of a given size?’

✿ Good lower bounds on $I_{\pi,P}$ translate into mixing time bounds for P .

$$T_*(\varepsilon \asymp 1) \lesssim \int_{\chi^2(\mu_0, \pi)^{-1}}^{1/2} \frac{p \, dp}{I_{\pi,P}(p)^2}.$$

✿ I will not go into the technical details of how this is achieved today.

► (...but \exists bonus slides).

Estimating the profile $I_{\pi,P}$

- ✂ Directly computing $I_{\pi,P}$ involves a difficult infimum over measurable sets.
- ✂ Our route will be to show that under verifiable conditions, we can estimate $I_{\pi,P}$.
- ✂ These conditions are nicely decoupled as \approx :
 1. A global condition about the target measure π .
 2. A local condition about the kernel P .
- ✂ Remark: This part of the analysis should work for \sim generic problems.

Isoperimetric Profiles of Probability Measures

- ✂ For $A \subseteq E$ and $r \geq 0$, let $A_r := \{x \in E : d(x, A) \leq r\}$.
- ✂ Define the *Minkowski content* of A under π with respect to d by

$$\pi^+(A) = \lim_{r \rightarrow 0^+} \inf \frac{\pi(A_r) - \pi(A)}{r}.$$

► \approx ‘boundary mass’ of A under π

- ✂ The *isoperimetric profile* of π with respect to the metric d is

$$I_\pi(p) := \inf \{ \pi^+(A) : A \subseteq E, \pi(A) = p \}, \quad p \in (0, 1).$$

- ✂ (usually) increasing on $[0, \frac{1}{2}]$, symmetric about $1/2$.
- ✂ For experts: This is (basically) $I_{\pi, P}$ for the Langevin diffusion.

Isoperimetric Profiles: Interpretation

- ✂ Isoperimetry relates the mass of sets to the mass of their boundaries.
- ✂ For Markov chains: isoperimetry captures how difficult it is to escape a given set.
- ✂ Escaping small sets ($p \rightarrow 0^+$) happens to be the relevant limit.
- ✂ If you escape all sets equally easily ($I_\pi(p) \geq c \cdot p$),
 - ▶ then you mix exponentially quickly.
- ✂ If you also escape small sets particularly well ($I_\pi(p) \gg c \cdot p$),
 - ▶ then things can be even better at the start.
- ✂ If small sets are hard to escape ($I_\pi(p) \ll c \cdot p$),
 - ▶ then things can be much worse.

Isoperimetric Profiles: Examples

- ✿ $\pi = \mathcal{N}(0, I_d)$ has $I_\pi(p) = (\varphi_\gamma \circ \Phi_\gamma^{-1})(p) \sim p \cdot \left(2 \cdot \log \frac{1}{p}\right)^{1/2}$ as $p \rightarrow 0^+$.
- ✿ $\pi(dx) \propto \exp(-|x|) dx$ has $I_\pi(p) = \min\{p, 1-p\}$.
- ✿ $\pi(dx) \propto \exp(-|x|^\alpha) dx$ has $I_\pi(p) \geq K(\alpha) \cdot p \cdot \left(\log \frac{1}{p}\right)^{1-1/\alpha}$ for $p \in [0, \frac{1}{2}]$.
- ✿ For log-concave measures,
 - ▶ \approx preserved under products.
 - ▶ functional inequalities (PI, LSI, \dots) imply bounds on I_π .
- ✿ Profiles transfer nicely under Lipschitz transport, bounded change of measure.
- ✿ Can be hard to obtain good bounds in some cases.
- ✿ Typically very informative.

‘Close Coupling’ of Markov Kernels

- ✿ Say that P is (d, δ, τ) -close coupling if for some **fixed** $\delta, \tau > 0$, it holds that

$$d(x, y) \leq \delta \implies \text{TV}(P_x, P_y) \leq 1 - \tau.$$

- ✿ When two chains get close enough, anywhere in the space,
 - ▶ there is a decent chance to make them coalesce.
- ✿ In our experience,
 - ▶ weaker assumption than e.g. global contractivity of dynamics ,
 - ▶ typically holds with better constants than minorisation conditions.
- ✿ δ is often small (but not tiny).
- ✿ τ can be of constant order (e.g. $1/4$).
- ✿ Remark: this condition can hold, with good (δ, τ) , for chains which mix **badly**.

Obtaining $I_{\pi,P}$

✿ Suppose that π has profile I_π , and P is (d, δ, τ) -close coupling. Then

$$I_{\pi,P}(p) \gtrsim \tau \cdot \min\{p, \delta \cdot I_\pi(p)\}$$

✿ Interpretation:

- ▶ If P is 'nice' at small scales,
- ▶ and if π is 'nice' at large scales,
- ▶ then P will mix well!

✿ For algorithms: no point in making τ too big; think of it as a constant.

- ▶ \rightsquigarrow Tune algorithm to find a good δ which gives a desired τ .

Isoperimetry: from π to P , to mixing

✂ Corollary 1: L^2 mixing time satisfies

$$T_*(\varepsilon \asymp 1) \lesssim \tau^{-2} \cdot \delta^{-2} \cdot \int_{\chi^2(\mu_0, \pi)^{-1}}^{1/2} \frac{p \, dp}{I_\pi(p)^2}.$$

(overlooking an additional annoying term related to the min)

✂ Corollary 2: for log-concave π , it holds that

$$\gamma_P \gtrsim \tau^2 \cdot \delta^2 \cdot I_\pi \left(\frac{1}{2} \right)^2.$$

✂ Our target is fixed, now: look at the kernel P , and control (τ, δ) .

Close Coupling for RWM

- ✂ For MH algorithms, natural to try

$$\mathrm{TV}(P_x, P_y) \leq \mathrm{TV}(P_x, Q_x) + \mathrm{TV}(Q_x, Q_y) + \mathrm{TV}(Q_y, P_y).$$

This appears to have some limitations.

- ▶ Roughly: tail behaviour ruins two of the three terms.
- ✂ Being ‘Metropolis-type’ (not just ‘Metropolis-Hastings-type’) lets us do better.
 - ▶ $\alpha(x, x') = \text{Monotone}(f(x')/f(x))$.
 - ▶ No ‘cross terms’, as in general MH.
- ✂ We will see that it suffices for us to control the acceptance rates.
 - ▶ \rightsquigarrow need to control the regularity of π .

Total Variation Bound between Metropolis Kernels

- ✿ Lemma: Let P be a Metropolis kernel, and suppose that $\inf_{x \in E} \alpha(x) \geq \alpha_0 > 0$. Then for any $x, y \in E$, it holds that

$$\text{TV}(P_x, P_y) \leq \text{TV}(Q_x, Q_y) + (1 - \alpha_0).$$

- ✿ Proof: Explicitly construct a coupling (next slide).

Proof Sketch

- ✂ WLOG, assume that $\pi(x) \geq \pi(y)$.
- ✂ If both chains propose moving to z , then $\alpha(x, z) \leq \alpha(y, z)$.
- ✂ Thus, can couple the acceptance steps so that almost surely,

$$x \text{ accepts move} \implies y \text{ accepts move}$$

- ✂ Use $P(A \cap B) \geq P(A) + P(B) - 1$ to see that

$$\begin{aligned} P(X' = Y') &\geq P(\tilde{X} = \tilde{Y}) + P(X' = \tilde{X}) - 1 \\ &\geq (1 - \text{TV}(Q_x, Q_y)) + \alpha_0 - 1 \\ &= \alpha_0 - \text{TV}(Q_x, Q_y). \end{aligned}$$

- ✂ Conclude by coupling inequality.

Acceptance Rate Bounds for RWM

- ✂ Recall that $\alpha(x, x') = \min \left\{ 1, \frac{\pi(x')}{\pi(x)} \right\}$.
- ✂ Natural to control growth of $U = -\log \pi$.
- ✂ Assumption: for some ψ , it holds that

$$U(x+h) - U(x) - \langle \nabla U(x), h \rangle \leq \psi(|h|).$$

- ✂ Lemma: The acceptance rate satisfies

$$\alpha(x) \geq \frac{1}{2} \cdot \exp \left(- \int \mathcal{N}(dz; 0, I_d) \cdot \psi(\sigma \cdot |z|) \right),$$

and taking $\sigma = v \cdot d^{-1/2}$ gives that

$$\alpha(x) \geq \frac{1}{2} \cdot \exp \left(-\psi(v) + \mathcal{O}(d^{-1}) \right).$$

Close Coupling for RWM

✂ Taking $\sigma = v \cdot d^{-1/2}$ allows for $\alpha_0 \geq \frac{1}{2} \cdot \exp(-\psi(v) + \mathcal{O}(d^{-1}))$.

✂ Taking $\delta = \sigma \cdot \alpha_0$ allows for

$$d(x, y) \leq \delta \implies \text{TV}(Q_x, Q_y) \leq \frac{1}{2} \cdot \alpha_0.$$

► (compute KL between Gaussians; apply Csiszar-Kullback-Pinsker)

✂ Using the coupling result,

$$\begin{aligned} \text{TV}(P_x, P_y) &\leq \text{TV}(Q_x, Q_y) + (1 - \alpha_0) \\ &\leq 1 - \frac{1}{2} \alpha_0, \end{aligned}$$

i.e. one may take $\tau = \frac{1}{2} \cdot \alpha_0$.

Isoperimetric Profile and Mixing of RWM

✿ Recalling that

$$I_{\pi,P}(p) \gtrsim \tau \cdot \min \{p, \delta \cdot I_{\pi}(p)\}$$

and taking ν so that $\alpha_0 \asymp 1$, obtain that

$$I_{\pi,P}(p) \gtrsim \min \{p, \sigma \cdot I_{\pi}(p)\},$$

$$\gamma_P \gtrsim \sigma^2 \cdot I_{\pi} \left(\frac{1}{2} \right)^2$$

$$T_*(\varepsilon \asymp 1) \lesssim \sigma^{-2} \cdot \int_{\chi^2(\mu_0, \pi)^{-1}}^{1/2} \frac{p \, dp}{I_{\pi}(p)^2}.$$

✿ Still very general at this stage.

Deducing main results (1)

- ✂ Under m -strong log-concavity, can bound isoperimetric profile as

$$I_{\pi}(p) \geq c \cdot m^{1/2} \cdot p \cdot \left(\log \frac{1}{p} \right)^{1/2}$$

- ✂ Under L -smoothness, take $\sigma = v \cdot (L \cdot d)^{-1/2}$ and control acceptance ratio as

$$\alpha_0 \geq \frac{1}{2} \cdot \exp \left(-\frac{1}{2} v^2 \right).$$

- ✂ Good isoperimetry, good acceptance rates \rightsquigarrow Good mixing.

► Remark: Failure of these conditions corresponds to known failure modes for RWM.

Deducing main results (2)

✎ Combining earlier results, obtain

$$\begin{aligned}\gamma_P &\gtrsim 1/(\kappa \cdot d) \\ T_*(\varepsilon \asymp 1) &\lesssim \sigma^{-2} \cdot m^{-1} \cdot \int_{\chi^2(\mu_0, \pi)^{-1}}^{1/2} \frac{dp}{p \cdot \log\left(\frac{1}{p}\right)} \\ &\lesssim \kappa \cdot d \cdot \log \log \chi^2(\mu_0, \pi).\end{aligned}$$

✎ Same strategy works well for other targets:

- ▶ Characterise the isoperimetric profile (out of your hands).
- ▶ Control the acceptance rates.

Not discussed in detail

- ✿ Sharpness of bounds w.r.t. d .
- ✿ Implications for asymptotic variance.
- ✿ ‘Multi-phase convergence’, initialisation.
- ✿ RWM on targets ‘between exponential and Gaussian’.
- ✿ RWM on rougher targets.
- ✿ pCN for Gaussian prior, ‘centered’ log-concave likelihood.

Ongoing and future work

- ✿ RWM on Heavy-tailed targets.
- ✿ Other Metropolis algorithms.
- ✿ Other non-Metropolis algorithms.
- ✿ New algorithms inspired by proof techniques.

Recap

- ✂ RWM for MCMC sampling.
- ✂ MCMC Convergence analysis via:
 - ▶ Isoperimetry (of target), and
 - ▶ Close Coupling (of kernels).
- ✂ Explicit control of RWM acceptance rates.
- ✂ Estimates of spectral gap, L^2 mixing times, asymptotic variance, etc.

Bonus Slides 1 : Technical Details

✿ Isoperimetric Profile to Conductance Profile:

$$\Phi_P(v) := \inf \left\{ \frac{I_{\pi, P}(v)}{v} : 0 < v \leq \frac{1}{2} \right\}. \quad (1)$$

✿ Conductance Profile to Spectral Profile:

$$\Lambda_P(v) \geq \frac{1}{2} \Phi_P(v)^2. \quad (2)$$

✿ Spectral Profile to Functional Inequality: for $f \geq 0$,

$$\frac{\mathcal{E}(P, f)}{\text{Var}_{\pi}(f)} \geq \frac{1}{2} \cdot \Lambda_P \left(4 \cdot \frac{\pi(f)^2}{\text{Var}_{\pi}(f)} \right). \quad (3)$$

✿ Functional Inequality to Mixing Time: consider $\|P^n f\|_2^2$ with $f = \frac{d\mu}{d\pi}$.

Bonus Slides 2 : Super-Poincaré Inequalities

✂ Spectral Profile to Super-Poincaré Inequality

$$\mathrm{Var}_{\pi}(f) \leq s \cdot \mathcal{E}(P, f) + \beta_P(s) \cdot \pi(|f|)^2; \quad (4)$$

can express β_P in terms of Λ_P .